

COMP9313: Big Data Management

MapReduce

w1: 10	goal
w2: 9	10: 2
w3: 10	9: 1
<w, fre>	<fre, count>

Data Structure in MapReduce

键值对

- Key-value pairs are the basic data structure in MapReduce

整数型

原始字节

- Keys and values can be: integers, float, strings, raw bytes
- They can also be arbitrary data structures

任意数据结构

<key, value>, key指的是word, value指的是frequency, 出现的次数

- The design of MapReduce algorithms involves: 包括将键值结构强加到任意数据集上

- Imposing the key-value structure on arbitrary datasets

- E.g., for a collection of Web pages, input keys may be URLs and values may be the HTML content

- In some algorithms, input keys are not used (e.g. wordcount), in others they uniquely identify a record

他们唯一地标识一个记录

- Keys can be combined in complex ways to design various algorithms

结果是要<value, count>

Recall of Map and Reduce

- Map

- Reads data (split in Hadoop, RDD in Spark)
- Produces key-value pairs as intermediate outputs

产生键值对作为中间输出

- Reduce

- Receive key-value pairs from multiple map jobs
- aggregates the intermediate data tuples to the final output

将中间数据元组聚合到最终输出

MapReduce in Hadoop

- Data stored in HDFS (organized as blocks)
 - Hadoop MapReduce Divides input into fixed-size pieces, input splits 将输入分割成固定大小的块，将输入分割
 - Hadoop creates one map task for each split
 - Map task runs the user-defined map function for each record in the split
 - Size of a split is normally the size of a HDFS block
 - Data locality optimization 数据本地优化
 - Run the map task on a node where the input data resides in HDFS 在输入数据驻留在HDFS中的节点上运行map任务
 - This is the reason why the split size is the same as the block size
 - The largest size of the input that can be guaranteed to be stored on a single node
 - If the split spanned two blocks, it would be unlikely that any HDFS node stored both blocks
- 如果拆分跨越两个块，则任何HDFS节点都不可能同时存储两个块

MapReduce in Hadoop

- Map tasks write their output to local disk (not to HDFS)
中间产出
 - Map output is intermediate output
 - Once the job is complete the map output can be thrown away
同步复制；复写
 - Storing it in HDFS with replication, would be overkill
 - If the node of map task fails, Hadoop will automatically rerun the map task on another node
- Reduce tasks don't have the advantage of data locality
 - Input to a single reduce task is normally the output from all mappers
单个化简任务的输入通常是所有映射器的输出
 - Output of the reduce is stored in HDFS for reliability
可靠性
 - The number of reduce tasks is not governed by the size of the input, but is specified independently
化简任务的数量不受输入大小的限制，而是独立指定的

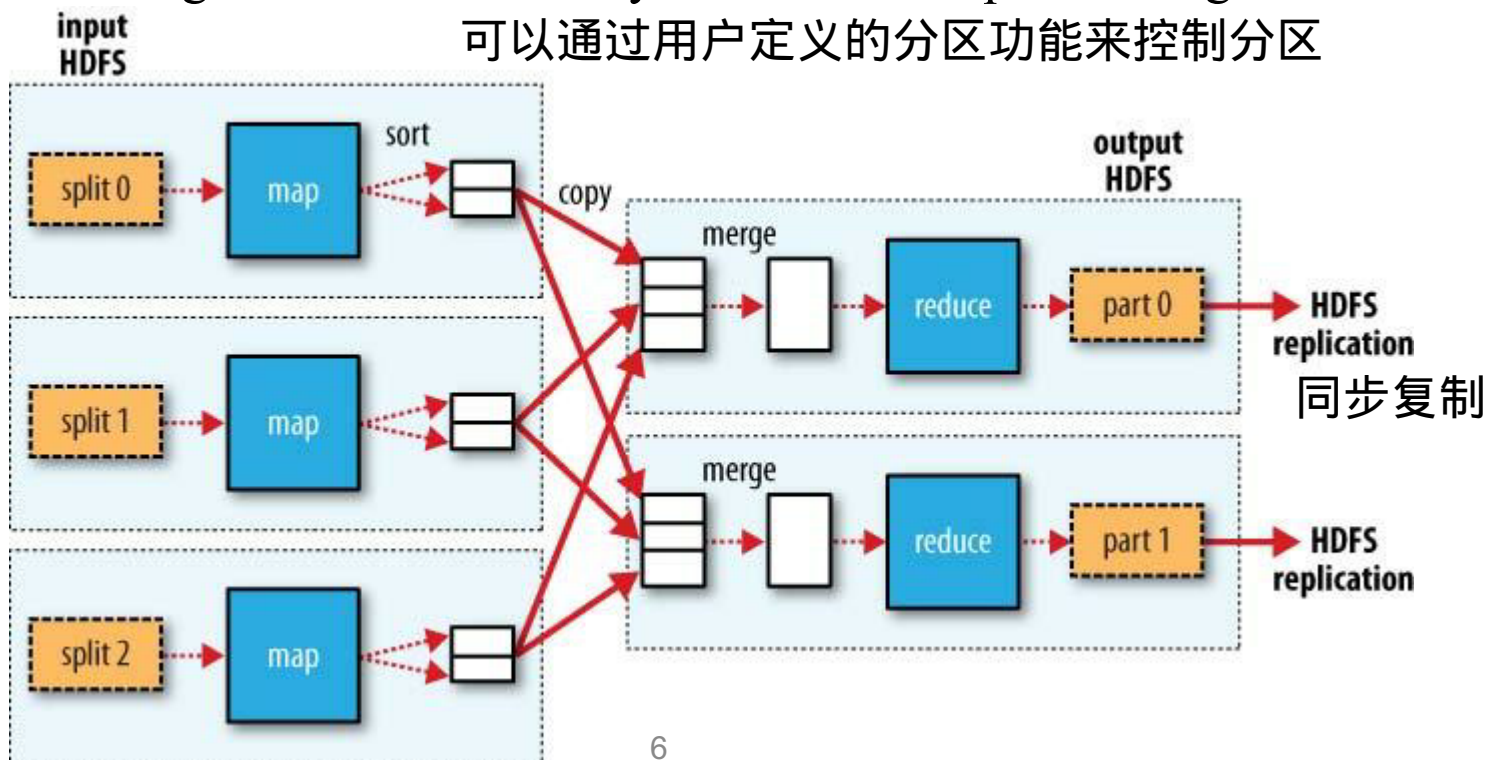
More Detailed MapReduce Dataflow

多个

分割

- When there are multiple reducers, the map tasks partition their output:
 - One partition for each reduce task
 - The records for every key are all in a single partition
 - Partitioning can be controlled by a user-defined partitioning function

可以通过用户定义的分区功能来控制分区



洗牌

Shuffle

洗牌是数据重新分配的过程

- Shuffling is the process of data redistribution
 - To make sure each reducer obtains all values associated with the same key.
 - It is needed for all of the operations which require grouping
 - E.g., word count, compute avg. score for each department, ...
- Spark and Hadoop have different approaches implemented for handling the shuffles.

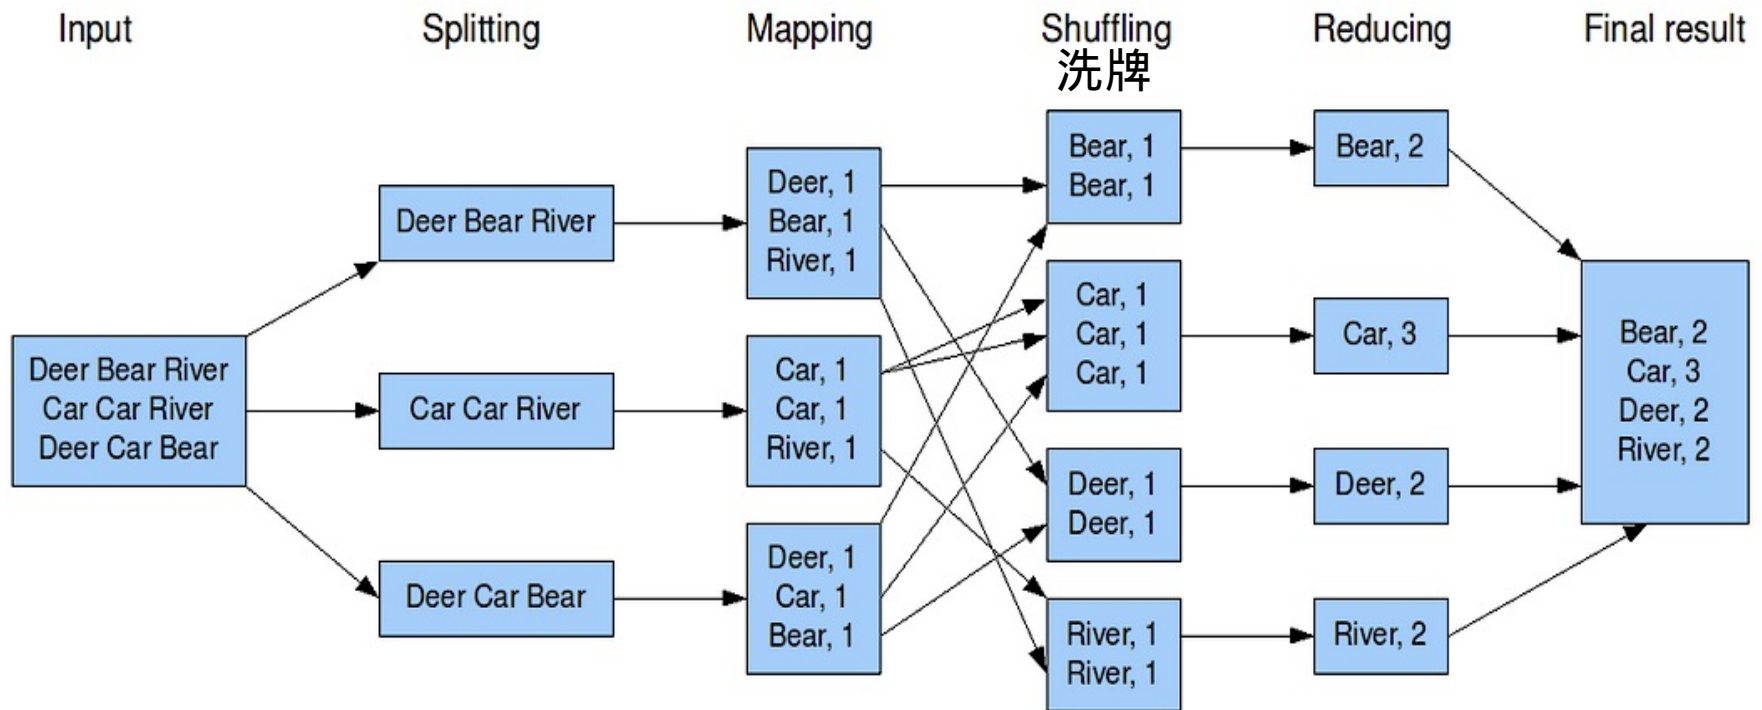
处理洗牌

Shuffle in Hadoop (handled by framework)

- Happens between each Map and Reduce phase 阶段
- Use Shuffle and Sort mechanism 使用洗牌和排序机制
 - Results of each Mapper are sorted by the key
 - Starts as soon as each mapper finishes
- Use combiner to reduce the amount of data shuffled 合成器
 - Combiner combines key-value pairs with the same key in each par 相同的键
 - This is not handled by framework!

Example of MapReduce in Hadoop

The overall MapReduce word count process



Shuffle in Spark (handled by Spark)

由某些操作触发

- Triggered by some operations
 - Distinct, join, repartition, all *By, *ByKey
 - 不同, 加入, 重新分区
 - I.e., Happens between stages
- Hash shuffle
- Sort shuffle
- Tungsten shuffle-sort
 - More on <https://issues.apache.org/jira/browse/SPARK-7081>

Hash Shuffle

分区

- Data are hash partitioned on the map side
 - Hashing is much faster than sorting
- Files created to store the partitioned data portion
 - # of mappers X # of reducers
- Use consolidateFiles to reduce the # of files
 - From $M * R \Rightarrow E * C/T * R$
- Pros:
 - Fast
 - No memory overhead of sorting
- Cons:
 - Large amount of output files (when # partition is big)

Sort Shuffle

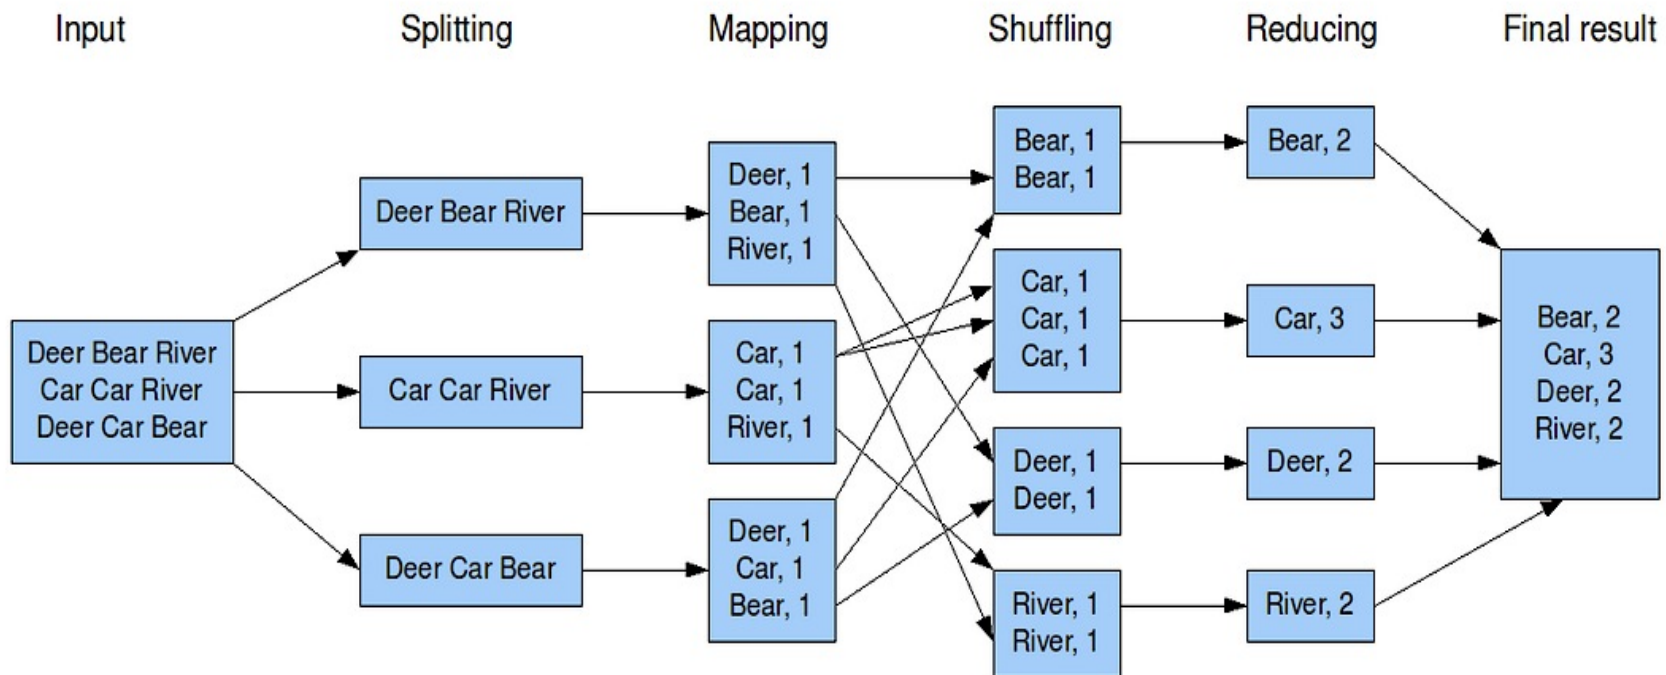
- For each mapper 2 files are created
 - Ordered (by key) data
 - Index of beginning and ending of each 'chunk'
- Merged on the fly while being read by reducers
- Default way
 - Fallback to hash shuffle if # partitions is small
- Pros
 - Smaller amount of files created
- Cons
 - Sorting is slower than hashing

MapReduce Functions in Spark (Recall)

- Transformation
 - Narrow transformation
 - Wide transformation
 - Action
-
- The job is a list of Transformations followed by one Action
 - Only action will trigger the 'real' execution 执行 ; 实施
 - I.e., lazy evaluation

Transformation = Map? Action = Reduce?

The overall MapReduce word count process



combineByKey

- $\text{RDD}([K, V])$ to $\text{RDD}([K, C])$
 - K: key, V: value, C: combined type
- Three parameters (functions)
 - createCombiner
 - What is done to a single row when it is FIRST met?
 - $V \Rightarrow C$
合并
 - mergeValue
 - What is done to a single row when it meets a previously reduced row?
 - $C, V \Rightarrow C$
 - In a partition
 - $C \Rightarrow C$
合并
 - mergeCombiners
 - What is done to two previously reduced rows?
 - $C, C \Rightarrow C$
 - Across partitions

Example: word count

- **createCombiner**
 - What is done to a single row when it is FIRST met?
 - $V \Rightarrow C$
 - $\text{lambda } v: v$
- **mergeValue**
 - What is done to a single row when it meets a previously reduced row?
 - $C, V \Rightarrow C$
 - $\text{lambda } c, v: c+v$
- **mergeCombiners**
 - What is done to two previously reduced rows?
 - $C, C \Rightarrow C$
 - $\text{lambda } c1, c2: c1+c2$

Example 2: Compute Max by Keys

- **createCombiner**
 - What is done to a single row when it is FIRST met?
 - $V \Rightarrow C$
 - $\lambda v: v$
- **mergeValue**
 - What is done to a single row when it meets a previously reduced row?
 - $C, V \Rightarrow C$
 - $\lambda c, v: \max(c, v)$
- **mergeCombiners**
 - What is done to two previously reduced rows?
 - $C, C \Rightarrow C$
 - $\lambda c1, c2: \max(c1, c2)$

Example 3: Compute Sum and Count

- `createCombiner`
 - $V \Rightarrow C$
 - `lambda v: (v, 1)`
- `mergeValue`
 - $C, V \Rightarrow C$
 - `lambda c, v: (c[0] + v, c[1] + 1)`
- `mergeCombiners`
 - $C, C \Rightarrow C$
 - `lambda c1, c2: (c1[0] + c2[0], c1[1] + c2[1])`

Example 3: Compute Sum and Count

- data = [('A', 2.), ('A', 4.), ('A', 9.), ('B', 10.), ('B', 20.), ('Z', 3.), ('Z', 5.), ('Z', 8.)]
 - Partition 1: ('A', 2.), ('A', 4.), ('A', 9.), ('B', 10.)
 - Partition 2: ('B', 20.), ('Z', 3.), ('Z', 5.), ('Z', 8.)
- Partition 1 ('A', 2.), ('A', 4.), ('A', 9.), ('B', 10.)
 - A=2. --> createCombiner(2.) ==> accumulator[A] = (2., 1)
 - A=4. --> mergeValue(accumulator[A], 4.) ==> accumulator[A] = (2. + 4., 1 + 1) = (6., 2)
 - A=9. --> mergeValue(accumulator[A], 9.) ==> accumulator[A] = (6. + 9., 2 + 1) = (15., 3)
 - B=10. --> createCombiner(10.) ==> accumulator[B] = (10., 1)
- Partition 2 ('B', 20.), ('Z', 3.), ('Z', 5.), ('Z', 8.), ('Z', 12.)
 - B=20. --> createCombiner(20.) ==> accumulator[B] = (20., 1)
 - Z=3. --> createCombiner(3.) ==> accumulator[Z] = (3., 1)
 - Z=5. --> mergeValue(accumulator[Z], 5.) ==> accumulator[Z] = (3. + 5., 1 + 1) = (8., 2)
 - Z=8. --> mergeValue(accumulator[Z], 8.) ==> accumulator[Z] = (8. + 8., 2 + 1) = (16., 3)
- Merge partitions together
 - A ==> (15., 3)
 - B ==> mergeCombiner((10., 1), (20., 1)) ==> (10. + 20., 1 + 1) = (30., 2)
 - Z ==> (16., 3)
- Collect
 - ([A, (15., 3)], [B, (30., 2)], [Z, (16., 3)])


reduceByKey

- `reduceByKey(func)`
 - Merge the values for each key using `func`
 - E.g., `reduceByKey(lambda x, y: x + y)`
- `createCombiner`
 - `lambda v: v`
- `mergeValue`
 - `func`
- `mergeCombiners`
 - `func`

groupByKey

- groupByKey()


- Group the values for each key in the RDD into a single sequence.
- Data shuffle according to the key value in another RDD



reduceByKey

- Combines before shuffling
- Avoid using groupByKey

ReduceByKey



The Efficiency of MapReduce in Spark

- Number of transformations
 - Each transformation involves a linearly scan of the dataset (RDD)
- Size of transformations
 - Smaller input size => less cost on linearly scan
- Shuffles
 - data transferring between partitions is costly
 - especially in a cluster!
 - Disk I/O
 - Data serialization and deserialization
 - Network I/O

Number of Transformations (and Shuffles)

```
rdd = sc.parallelize(data)
```

- data: (id, score) pairs

- Bad design

```
maxByKey = rdd.combineByKey(...)
```

```
sumByKey = rdd.combineByKey(...)
```

```
sumMaxRdd = maxByKey.join(sumByKey)
```

- Good design

```
sumMaxRdd = rdd.combineByKey(...)
```


Size of Transformations

```
rdd = sc.parallelize(data)
```

- data: (word, 1) pairs

- Bad design

```
countRdd = rdd.reduceByKey(...)
```

```
fileteredRdd = countRdd.filter(...)
```

- Good design

```
fileteredRdd = countRdd.filter(...)
```

```
countRdd = fileteredRdd.reduceByKey(...)
```

Partition

```
rdd = sc.parallelize(data)
```

- data: (word, 1) pairs

- **Bad design**

```
countRdd = rdd.reduceByKey(...)
```

```
countBy2ndCharRdd = countRdd.map(...).reduceByKey(...)
```

- **Good design**

```
partitionedRdd = data.partitionBy(...)
```

```
countBy2ndCharRdd = partitionedRdd.map(...).reduceByKey(...)
```